

MATH 3190 Homework 8

Focus: Notes 10 and Notes 11

Due April 25, 2024

Your homework should be completed in R Markdown or Quarto and Knitted to an html or pdf document. You will “turn in” this homework by uploading to your GitHub Math_3190_Assignment repository in the Homework directory.

Problem 1 - Time Series (34 points)

In the `itsmr` package is the dataset called `wine` that contains 142 accounts of Australian red wine sales per month from January 1980 to October 1991.

Part a (3 points)

Load in this data and convert it to a `tsibble`. You can do this by first making it a `ts` object by typing `ts(wine, deltat = 1/12, start = 1980)` and then piping that into the `as_tsibble()` function in the `tsibble` package.

Then use `ggplot()` to plot the time series. Show both the points and lines connecting them. On the x axis, put each year using the `scale_x_yearmonth()` function and give the plot a title.

Part b (1 point)

Remove the last 34 rows from the wine data. We will use them to compare how our forecasts look later.

Part c (3 points)

Use the `gg_season()` plot in the `feasts` library to plot the seasonality. Does it seem like there is trend and/or seasonality in this time series? If so, explain the patterns.

Part d (3 points)

Use the `adf.test()` function in the `tseries` library to perform a Dickey-Fuller test to determine if we have evidence this time series is stationary. Only the column with the sales should be entered. Since we have data by month, use the `k = 12` option. Comment on the results of the test.

Part e (4 points)

Difference the data once. Take $Y_t = Y_{t+1} - Y_t$ for $t = 1, \dots, n - 1$. In the notes, I defined the differences as $Y_t = Y_t - Y_{t+1}$, but it actually does not matter. **Plot** this differenced time series (you can just put the index on the x-axis) and use the `adf.test()` function again on the differenced data. Again, use `k = 12` option. **Comment** on the results of the test.

Part f (3 points)

Now use the `auto.arima()` function in the `forecast` library to fit an ARIMA model to this time series (put the entire `tsibble` into this function).

Feel free to check out [this resource](#) that discusses more about time series, including what the drift term is.

Part g (4 points)

Use the `forecast()` function from the `forecast` library to predict the future for 34 months. You should specify you are using the `forecast` library by typing `forecast::forecast()` since the `itsmr` library also has a `forecast()` function. The input of this function is your model you fit in part f. Use the model fit to the undifferenced data.

Then use the `autoplot()` function to plot the forecast.

Part h (8 points)

Plot the actual values for 1990 and 1991 along with the predicted values from the forecast and a shaded 95% interval bands. The `geom_ribbon()` function may be useful here. Is there any point where the true value is outside the interval? If so, report the year(s) and month(s) where that is the case.

Part i (5 points)

Use the `checkresiduals()` function in the `forecast` library to create a residual plot, an autocorrelation plot, and a histogram for the residuals. The input to this function should be your model fit with the `auto.arima()` function. Comment on whether the assumptions seem to be satisfied here. If you'd like, you can make a QQ-plot as well, but the histogram contains the same information.

Note: this function will also give you information from the Ljung-Box test, which tests if the residuals exhibit autocorrelation. A large p-value means we have no evidence of autocorrelation.

Problem 2 (66 points)

The National Institute of Diabetes and Digestive and Kidney Diseases conducted a study on 768 adult female Pima Native Americans living near Phoenix. The purpose of the study was to investigate factors related to diabetes. The response variable in this study (`test`) was whether or not the patient showed signs of diabetes (1 = yes, 0 = no). One explanatory variable considered was the plasma glucose concentration (in mg/dL) at 2 hours in an oral glucose tolerance test (`glucose`). Many other potential explanatory variables were considered as well. The goal here is to predict the diabetes status of the subject.

Part a (5 points)

Install and/or load the `faraway` package with the `pima` dataset. Set a seed and then split the data into a testing and a training set with the testing set containing 20% of the data values. When you do this, change the name of the variable “test” to “result” since it will be a bit confusing to call that test when we have a test dataset. Then make the `result` variable a factor.

Part b (4 points)

Fit a logistic regression model to the training set for predicting `result` from all other variables. Then obtain the predictions for the `test` set with a probability cutoff of 0.5 to predict if the person has diabetes. Create a confusion matrix using the `confusionMatrix()` function in the `caret` library. Print all of the output from that function that shows the confusion matrix and many associated statistics.

Part c (10 points)

Now let's use naïve Bayes to get predictions. First, find the predicted probability of being a 0 and a 1 for the first observation in the training set “by hand” like we did in the notes. To simplify this a bit, let's only use the variables `pregnant` and `bmi` here. Of course, use `R` to help you (the `dnorm()` function and some tidyverse functions will be especially nice), but don't use the `naive_bayes()` function yet. Note: these variables we

are using to predict are numeric, so we will assume a normal distribution for both of them (probably not a good assumption, but let's go with it).

Part d (3 points)

Now use the `naive_bayes()` function in the `naivebayes` library to fit a model predicting `result` using `pregnant` and `bmi`. Find the predicted probabilities of being a 0 and a 1 for the first observation and show that it matches what you found in part c.

Part e (3 points)

Now fit a naïve Bayes model predicting `result` from **all** other variables. Then report the output of the `confusionMatrix()` function for this predictor when predicting on the **test** set.

Part f (3 points)

Now fit a naïve Bayes model predicting `result` from all other variables using a kernel instead of assuming a normal distribution. You can do this by setting `usekernel = TRUE` in the `naive_bayes()` function. Then report the output of the `confusionMatrix()` function for this predictor when predicting on the **test** set.

Part g (10 points)

Now let's use a support vector machine (SVM) for classifying diabetes using all other variables. We'll use the `svm()` function in the `e1071` library. Use a linear kernel, keep `scale = TRUE` (which is the default), and perform 5-fold cross validation to select the cost for costs between 0.1 and 5 with a step size of 0.1 (0.1, 0.2, 0.3, etc.). This will be similar to what you did back in Homework 3 problem 2h. Set a seed here. Like you did there, let's use Cohen's kappa to select the best cost value.

Part h (4 points)

Using the “best” cost value from part g, fit the SVM with a **linear** kernel and find the confusion matrix for predicting the **test** set.

Part i (10 points)

Now let's use a different kernel. Let's use the radial kernel for the SVM. This has two tuning parameters: the cost and γ . Let's use the `train()` function in the `caret` package to do 5-fold cross validation (set a seed, please) to select both of those parameters. However, the `train()` function uses σ instead of γ and C instead of cost.

Use the `trainControl()` function with the `trControl` option and the `tuneGrid` option in the `train()` function to perform this cross-validation. For `sigma`, use `seq(0.1, 2, by = 0.1)` and for `C` also use `seq(0.1, 2, by = 0.1)`. Put `method = "svmRadial"` in the `train()` function as well so it knows to use a SVM with the radial kernel. You will need to install the `kernlab` package for this to work.

By default, the “best” `sigma` and `C` values will be chosen with accuracy instead of Cohen's kappa. Make sure to choose the best values using kappa (although both metrics may choose the same “best” values).

Part j (4 points)

Fit the SVM with the **radial** kernel using the `svm()` function and the cross-validated cost and γ found in part i and print a confusion matrix for predicting on the **test** set.

Part k (10 points)

Compare the output of the `confusionMatrix()` function for all of our classifiers: the logistic regression model, the naïve Bayes with a normal assumption, the naïve Bayes using kernel density estimation, the SVM with a

linear kernel, and the SVM using a radial kernel. Which was best? Which was worst? How different were the results?